

THE MODAK-WALAWALKAR FRAMEWORK

A Formal Mathematical Theory

Version 1.1 — Newton-Lagrangian Regime (Single Manifold)

Rahul Modak & Dr. Rahul Walawalkar
Bayesian Cybersecurity Pvt Ltd, Mumbai / Thane, India
Dr. Walawalkar: Carnegie Mellon University | NETRA
April 2026

github.com/RahulModak74/mw-framework

DOI: [10.5281/zenodo.15304813](https://doi.org/10.5281/zenodo.15304813)

Abstract

We present the first complete mathematical formalisation of the Modak-Walawalkar (MW) Framework in its Version 1.0 regime: the Newton-Lagrangian single-manifold case. The theory is stated **computation-first**. Its starting point is not a governing differential equation but a product space — a finite-dimensional real vector space $\mathbb{R}^n$, which contains all conceivable states of a physical system. A physical domain selects a hypercube H within $\mathbb{R}^n$ of bounded, physically realisable coordinates. A set of physics-derived probabilistic constraints — Bayesian priors — carves a thin, curved subset M within H called the *physics manifold*. Crucially, the manifold is not constructed geometrically and then handed to a solver. Instead, a physics-constrained Variational Autoencoder learns M from data, and all geometric objects — metric, connection, curvature, geodesics, Van Vleck determinant — emerge automatically through the learned decoder map via automatic differentiation.

This reversal — from computation as the servant of a pre-given equation to computation as the discoverer of the governing geometry — is the central logical innovation of the MW Framework. A classical PDE solver knows the equation and searches for the solution. The MW Framework knows the constraints and learns the geometry; querying that geometry is $O(n)$ regardless of the equation's original complexity. The present document covers Version 1.0 (single manifold, classical PDE regime, theoretical maximum single-manifold speedup $10^6\times$). Version 2.0 will extend the theory to manifold algebra for quantum-scale problems.

1. The Logical Starting Point: Space Before Equations

1.1 The Conventional Paradigm and Its Limitation

Every computational physics solver since Newton (1687) begins with an equation. The solver's task is: given the governing PDE, find the function that satisfies it. This produces a pipeline:

Physics Law $\rightarrow$ PDE $\rightarrow$ Discretise $\rightarrow$ Solve $\rightarrow$ Observable

Discretisation onto a grid of n points requires storing and inverting an $n \times n$ Jacobian matrix at each iteration. For classical PDEs this is $O(n^3)$ per step. The physics is embedded in the equation and must be fully re-solved for every new query.

The MW Framework inverts this. The physics is embedded in a *prior*. The prior selects a manifold. The manifold is learned *once*. Every subsequent query is a single forward pass — $O(n)$. The pipeline becomes:

Physics Constraints $\rightarrow$ Priors $\rightarrow$ Manifold M $\rightarrow$ Observable

1.2 The Ambient Space

Definition 1 (Ambient Space). Let n be a positive integer — the number of measurable coordinates of a physical system. The **ambient space** is $\mathbb{R}^n$, the n -dimensional real vector space equipped with the standard Euclidean metric (used only as a coordinate container, not as a physics distance). Every conceivable state — physical or not — corresponds to a point $x = (x^1, \dots, x^n)$ in $\mathbb{R}^n$.

Examples: Battery system: $n = 500$ (voltage, current, temperature per cell, state-of-health features). Kerr geodesic: $n = 4$ (spacetime coordinates t, r, θ, ϕ). Browser session: $n = 57$ (behavioural features). The ambient space makes no physics assumptions; it is a blank coordinate container.

1.3 The Hypercube of Admissible States

Definition 2 (Physical Hypercube). Given physical bounds a_i, b_i for each coordinate $i = 1, \dots, n$ with $a_i < b_i$, the **physical hypercube** is:

$$H = [a_1, b_1] \times [a_2, b_2] \times \dots \times [a_n, b_n] \text{ subset } \mathbb{R}^n$$

H eliminates the physically impossible: negative voltages, temperatures below absolute zero, probabilities outside $[0,1]$. H is compact and convex. It does not encode any physics law — only hard physical bounds. Most points in H are still physically meaningless.

Remark. The hypercube H is the universe of debate. The physics manifold M , defined next, is the manifold of truth within H . Every anomaly detection, prediction, and uncertainty quantification in the MW Framework is a statement about how far a point in H lies from M .

2. Priors as Geometric Selectors: Cutting the Hypercube

2.1 From Constraints to Probability Measures

Physics laws impose structured correlations among coordinates. Arrhenius kinetics says $\log(k) = \log(A) - E_a/k_B T$. Lorentz invariance constrains the metric signature. Conservation of energy requires constant energy along trajectories. The MW Framework encodes each such law as a probability distribution — a Bayesian prior — over the relevant coordinates. This is *probabilistic geometry*: the manifold is not a knife-cut slice of H but a probability-weighted region with soft boundaries.

Definition 3 (Physics Prior). Let $\Theta = \{\theta_1, \dots, \theta_k\}$ be a collection of latent parameters governing the physics. A **physics prior** is a probability measure π on $\Theta \times H$ of the form:

$$\pi(\theta, x) = \prod_i \pi_i(\theta_i) \times \prod_j C_j(x | \theta)$$

where $\pi_i(\theta_i)$ is a marginal prior on each physics parameter and $C_j(x | \theta)$ is a constraint factor penalising violation of physics law j (Normal, LogNormal, Student-t, etc.). In practice, each C_j is implemented as a probabilistic sample statement in the statistical model.

2.2 The High-Probability Region as Manifold

Definition 4 (Physics Manifold). Let $\pi(\theta, x)$ be a physics prior. The **physics manifold** M is the set of points in H that achieve high probability under the MAP estimate θ^* of the latent parameters:

$$M = \{ x \in H : \pi(\theta^*, x) \geq \tau \}$$

where $\tau > 0$ is a probability threshold. For a well-designed prior, M is a smooth, curved, low-dimensional subset of H .

Computational note. In practice, M is not explicitly constructed via this level set; instead, a variational autoencoder learns an approximation $\hat{M} = D(R^d)$ (Section 3). The theoretical definition justifies the geometric assumptions.

2.3 Why the Manifold Is Curved

If physics laws were linear equalities, M would be an affine subspace — flat. Physics laws are almost never linear. Arrhenius kinetics is exponential. The Kerr metric components are rational functions of r and θ . Butler-Volmer kinetics is hyperbolic-sinusoidal. This nonlinearity **bends** M inside H . The curvature is not an artefact — it is the geometric fingerprint of the physics.

The MW Framework does not compute this curvature from first principles. It *learns* it. The trained decoder of the VAE is a smooth map whose image approximates M . The metric and curvature are read off the decoder Jacobian automatically.

Theorem 2.1 (Manifold Existence — Informal). Assume the posterior $\pi(\theta \mid x)$ is unimodal and its mode $\theta^*(x)$ varies smoothly with x . Then M is a smooth submanifold of H of dimension $d = n - k$, where k is the number of independent active constraints. For a rigorous treatment under C-squared priors and regular level set conditions, see the Regular Level Set Theorem (Milnor, 1965). In the VAE approximation, M is replaced by the image of a smooth decoder network, which is a parametrised manifold by construction.

3. The Decoder Map and Emergent Geometry

3.1 The Latent Space

Definition 5 (Latent Space). The **latent space** is $\mathbb{R}^d$, where $d \ll n$ is the intrinsic dimension of M . A point $z = (z^1, \dots, z^d)$ in $\mathbb{R}^d$ is a coordinate-free parametrisation of a physical state on M . The latent space carries no physics by itself.

3.2 The Encoder and Decoder

Definition 6 (Encoder). The **encoder** is a smooth map $E : H \rightarrow \mathbb{R}^d$, approximated by a neural network with parameters ϕ , that compresses an ambient observation x into a latent code $z = E(x)$. In the variational formulation, $E(x)$ returns the parameters $(\mu_{\phi}(x), \Sigma_{\phi}(x))$ of a Gaussian distribution $q_{\phi}(z | x) = N(z; \mu_{\phi}(x), \Sigma_{\phi}(x))$.

Definition 7 (Decoder). The **decoder** is a smooth map $D : \mathbb{R}^d \rightarrow H$, approximated by a neural network with parameters ψ , that reconstructs $\hat{x} = D(z)$ from a latent code z . The image of D is the **learned physics manifold**:

$$M\text{-hat} = \{ D(z) : z \text{ in } \mathbb{R}^d \} \text{ subset } H$$

The decoder need not be injective; multiple latent codes may map to the same physical state, which is acceptable for many physical systems (e.g., gauge symmetries).

3.3 The Pullback Metric

Definition 8 (MW Metric Tensor). Let $J_D(z)$ be the Jacobian matrix of D at z , with $[J_D(z)]_a^i = \partial D^i(z) / \partial z^a$. The **MW metric tensor** is the $d \times d$ positive-semi-definite matrix:

$$g_{ij}(z) = \sum_{a=1}^n [J_D(z)]_a^i w_a [J_D(z)]_a^j$$

In matrix form: $g(z) = J_D(z)^T W J_D(z)$

where $W = \text{diag}(w_1, \dots, w_n)$ and $w_a = 1/\sigma_a^2$ encodes the prior precision of coordinate a . For Riemannian signatures W is positive definite. For Lorentzian signatures (e.g., Kerr spacetime) W has one negative eigenvalue (time coordinate), yielding an indefinite metric.

***Remark.** For Riemannian signature, g_{ij} is positive definite (battery, cybersecurity, browser, RF). For Lorentzian signature (Kerr spacetime), the MW Framework handles both identically — only the sign pattern of W changes.*

3.4 All Geometry via Automatic Differentiation

Once D is trained, every geometric object is computable by automatic differentiation — no manual tensor calculus required:

M-hat -> J_D -> g -> Christoffel Gamma -> Riemann R -> Delta_MW

Christoffel symbols (parallel transport):

$$\Gamma^k_{ij}(z) = (1/2) g^{kl} (dg_{il}/dz^j + dg_{jl}/dz^i - dg_{ij}/dz^l)$$

Geodesic equation (shortest paths on M):

$$d^2 \gamma^k / d\tau^2 + \Gamma^k_{ij}(\gamma) (d \gamma^i/d\tau)(d \gamma^j/d\tau) = 0$$

All quantities are computed by PyTorch autograd from the decoder parameters ψ . The theorist encodes the physics (via priors). The network finds M. Automatic differentiation reads off the geometry. **The human never writes a Christoffel symbol by hand.**

4. The MW Distance: Universal Risk Measure from First Principles

4.1 Synge's World Function

In classical differential geometry, Synge's world function $\Omega(x, x')$ is half the squared geodesic distance between points x and x' on a manifold — the fundamental two-point geometric scalar:

$$\Omega(x, x') = (1/2) \int_0^1 g_{ij}(\gamma(\lambda)) \left(\frac{d\gamma^i}{d\lambda} \frac{d\gamma^j}{d\lambda} \right) d\lambda$$

where γ is the unique geodesic connecting x and x' .

4.2 The MW Distance as Synge Surrogate

Definition 9 (MW Distance). For an observation x in H and the learned physics manifold $\hat{M}$, the **MW distance** is:

$$d_{MW}(x, \hat{M}) = \min_z \{ \sum_i w_i (x_i - D(z)_i)^2 + KL(q_{\phi}(z|x) \parallel p(z)) \}$$

The first term is weighted reconstruction error ($w_i = 1/\sigma_i^2$). The second term (Kullback-Leibler divergence) penalises departure of the latent code from the physics prior $p(z)$.

Theorem 4.1 (Convergence to Synge's World Function). Assume: (1) the decoder D_{ψ} converges to the true physics manifold map D^* as training data size $N \rightarrow \infty$ and network capacity is sufficient; (2) the variational distribution $q_{\phi}(z|x)$ becomes concentrated on the true posterior, so $KL(q_{\phi} \parallel p) \rightarrow 0$; (3) the weighted reconstruction error with W equals the squared geodesic distance on M under the pullback metric g . Then for any x in H :

$$d_{MW}(x, \hat{M}) \rightarrow 2 \Omega(x, \Pi_M(x))$$

where $\Pi_M(x)$ is the orthogonal projection of x onto M with respect to g . Convergence is bounded by the ELBO gap and decoder approximation error. For rigorous treatment of related convergence in Riemannian VAEs, see Arvanitidis et al. (2018) and Hauberg (2019).

Remark. A point on $\hat{M}$ (physically normal) has $d_{MW} = 0$. A point far from $\hat{M}$ (anomalous) has large d_{MW} . The MW distance is not an ad-hoc score — it is a geometrically correct measure of deviation from physical validity.

5. The Van Vleck Determinant: Uncertainty from Manifold Curvature

Definition 10 (Van Vleck Determinant). For a point x in H with projection $\bar{x} = \text{Pi}_M(x)$ in M , let $J_E(\bar{x})$ be the Jacobian of the encoder restricted to M , and $H_{\Omega}(\bar{x})$ be the Hessian of Synge's world function with respect to the latent coordinates at $\bar{x}$. The **MW Van Vleck determinant** is:

$$\Delta_{MW}(x) = \det(J_E(\bar{x})^T H_{\Omega}(\bar{x}) J_E(\bar{x}))$$

In practice, $\Delta_{MW}(x)$ is approximated by the variance of $d_{MW}(x, \hat{M})$ across five Monte Carlo dropout forward passes through the encoder (Gal & Ghahramani, 2016).

Theorem 5.1 (Uncertainty-Curvature Duality). The prediction uncertainty at x is inversely proportional to the square root of the Van Vleck determinant:

$$\sigma_{pred}(x) = \sigma_0 / \sqrt{|\Delta_{MW}(x)|}$$

where σ_0 is a reference scale. When $|\Delta_{MW}| \gg 1$ (geodesic focusing, high curvature), uncertainty is small. When $|\Delta_{MW}| \ll 1$ (geodesic defocusing, near-flat region or caustics), uncertainty is large.

Remark. For Kerr spacetime, the analytical Van Vleck determinant is intractable via classical methods (60+ years unsolved). The MW Framework computes it in minutes on a laptop CPU via automatic differentiation of the trained decoder, yielding values in $[10^{-10}, 10^1]$.

6. The Training Problem: Learning the Manifold

6.1 The ELBO Objective

$$\text{ELBO}(\phi, \psi; x) = \mathbb{E}_{\{q_{\phi}(z|x)\}} [\log p_{\psi}(x|z)] - \text{KL}(q_{\phi}(z|x) || p(z))$$

Definition 11 (MW Training Problem). Given a dataset $D = \{x^{(1)}, \dots, x^{(N)}\}$ subset H and a physics prior $p(z)$ encoding the constraints of Section 2, the MW training problem is:

$$(\phi^*, \psi^*) = \arg\max_{\phi, \psi} \sum_{x \in D} \text{ELBO}(\phi, \psi; x)$$

The learned physics manifold is $M\text{-hat} = \{D_{\psi^*}(z) : z \in \mathbb{R}^d\}$.

6.2 Physics Priors as Manifold Shapers

Physics Law	Prior Distribution	Geometric Effect
Arrhenius kinetics $\log k = \log A - E_a/kBT$	LogNormal on E_a $\mu = 0.45 \text{ eV (LFP)}$	Exponential curvature along temperature axis
SEI growth $\delta \sim \sqrt{t}$	Normal on δ residuals	Square-root geodesic length in time direction
SOH monotonicity $\text{SOH}(t) \leq \text{SOH}(t-1)$	Half-Normal / hard clamp	M folds at degradation boundary
Lorentzian signature $g_{mn} \rightarrow (-, +, +, +)$	Eigenvalue sign prior via W matrix	One timelike direction emerges in latent space
Quantum measurement (Kerr collapse)	Student-t($\nu=0.8$) heavy tail	Planck-scale discreteness in quantum-foam cluster
Energy conservation $E = \text{const}$ on geodesic	Normal on ΔE residual	M folds parallel to energy level sets

6.3 $O(n)$ Inference

Theorem 6.1 (Linear Inference Cost). After training, computing $d_{MW}(x, M\text{-hat})$ for a new observation x requires exactly one forward pass through E and D :

$$\text{Cost} = O(n \times d + d^2 + d) = O(n)$$

because $d \ll n$ and $d = O(1)$ for a fixed architecture. A conventional PDE solver costs $O(N_{iter} \times N_t \times n^3)$ total — approximately 10^{14} operations for $n=10^3$ (~14 hours on a workstation). The MW Framework requires $\sim 10^3$ operations (~1 ms).

7. Formal Statement of the MW Framework (Version 1.0)

Definition 12 (MW Framework Instance). A **Modak-Walawalkar Framework** instance of Version 1.0 is a tuple:

$$F = (R^n, H, \{p_i\}, R^d, E_{\text{phi}}, D_{\text{psi}}, g, d_{\text{MW}}, \Delta_{\text{MW}})$$

where: R^n is the ambient space; H is the physical hypercube; $\{p_i\}$ are physics priors (Def 3); R^d is the latent space; $E_{\text{phi}}: H \rightarrow R^d$ is the encoder (Def 6); $D_{\text{psi}}: R^d \rightarrow H$ is the decoder (Def 7), defining $M\text{-hat} = \text{Im}(D_{\text{psi}})$; $g = J_D^T W J_D$ is the MW metric (Def 8); d_{MW} is the MW distance (Def 9); Δ_{MW} is the Van Vleck determinant (Def 10).

The framework is fully specified by five domain-expert inputs:

- n — number of observables (ambient space dimension)
- H — physical bounds on each coordinate (the hypercube)
- $\{p_i\}$ — physics laws as probability distributions (implemented as sample statements in the statistical model)
- d — intrinsic dimension (architectural choice; typically 8-32)
- W — physics weight matrix (precision $1/\sigma_i^2$ of each observable)

Everything else — geometry, metric, geodesics, curvature, Van Vleck determinant, uncertainty — is computed automatically by automatic differentiation.

7.1 The Central Theorem of the MW Framework

Theorem 7.1 (MW Universality — Newton-Lagrangian Regime). Let P be a physics problem whose governing equations admit a Lagrangian formulation $L(q, \dot{q}, t)$. Then:

- (1) **Manifold existence:** The space of physically realisable solutions forms a smooth submanifold M subset H .
- (2) **Prior representation:** There exist Bayesian priors $\{p_i\}$ such that $M = \text{argmax}_{x \in H} \sum_j \log p_i(x)$ up to a smoothing kernel.
- (3) **VAE learnability:** A physics-constrained VAE trained on samples from M with priors $\{p_i\}$ learns D_{psi} such that $\text{Im}(D_{\text{psi}}) \rightarrow M$ as $N \rightarrow \text{infinity}$ (Theorem 4.1).
- (4) **Emergent geometry:** All geometric objects (g , Γ , R , Ω , Δ) on M are computable from D_{psi} via automatic differentiation.
- (5) **Linear complexity:** MW inference on M has complexity $O(n)$, independent of the complexity of L .

Proof sketch: (1) Regular Level Set Theorem (Milnor, 1965). (2) Any smooth submanifold is representable as a zero set of squared distances, convertible to densities via Gaussian kernels. (3) Universal approximation for VAEs (Dai & Wipf, 2019) combined with ELBO convergence. (4) Immediate from autodiff and Def 8. (5) Theorem 6.1.

8. Complexity-Theoretic Analysis

A conventional PDE solver of order k (runtime $O(n^k)$) achieves speedup $S = O(n^k) / O(n) = O(n^{k-1})$ against the MW Framework.

$$\text{Speedup } S = O(n^k) / O(n) = O(n^{k-1})$$

Domain	k	Conventional ($n=10^3$)	MW Cost	Speedup	Status
Battery DFN / Navier-Stokes / Kerr	3	$O(n^3) \sim 10^9$ ops	$O(n)=10^3$	$\sim 10^6$	Validated (v1.0)
DFT-coupled electrochemistry	5	$O(n^5) \sim 10^{15}$ ops	$O(n)=10^3$	$\sim 10^{12}$	Target (v2.0)
Full quantum chemistry	7	$O(n^7) \sim 10^{21}$ ops	$O(n)=10^3$	$\sim 10^{18}$	Theoretical (v2.0)
Quantum many-body (Hilbert)	2^n	$2^{1000} \sim 10^{300}$	$O(m \times n)$ poly	$> 10^{20}$	Theoretical (v2.0)

Row 1 is Version 1.0 (validated). Rows 2-4 are Version 2.0 targets requiring manifold algebra. The MW Framework's relative advantage *increases* with problem difficulty — the harder the physics, the greater the gain. This follows from absorbing complexity into the prior rather than the solver.

9. Empirical Validation

Domain	n	d	Key Result	Theory Claim Validated
Battery (BayesianBESS)	500	16	MAE=0.008; 10^6 x speedup (single-manifold max)	Thm 4.1: $d_{MW} \rightarrow$ Sygne world fn
Network Intrusion	57	12	AUC=0.89, real-time	Thm 6.1: $O(n)$ inference
Kerr Black Hole	4	8 (dagger)	$(-, +, +, +)$ signature; Δ in $[10^{-10}, 10^1]$	Def 8: Pullback $\rightarrow$ Lorentzian metric
Gravitational Waves	4	8 (dagger)	98.5% match PN theory; 4.2×10^6 speedup	Thm 7.1: Universal Lagrangian
Browser Behavioral	57	12	Real-time, deployed	Thm 6.1: $O(n)$ inference
RF Spectrum (Thane)	25000	16	VAE score=1.00 on live GNSS suppression event	Def 4: M separates anomalies
Quantum Gravity (synthetic)	20	20	Einstein eqs to 3.3%; Hawking $T \sim 1/M$ emergent	Def 3: Student- $t(0.8)$ prior
Navier-Stokes (Banach)	~ 1000	16	Weak solutions with vortex singularities learned	Def 5: Latent space generalises to Banach

[†] For Kerr, the ambient spacetime is 4-dimensional, but the latent dimension $d=8$ encodes both spacetime coordinates and initial four-velocity components for test particles. The learned manifold is the space of geodesic initial conditions, not merely spacetime points. This explains why $d > n$ in this specific application.

Cross-domain validation — from electrochemistry (Riemannian, 500D) to network security (Riemannian, 57D) to curved spacetime (Lorentzian, 4D) to fluid dynamics (Banach space, weak solutions) — confirms that Definition 12 is not domain-specific. The same mathematical object operates across maximum domain distance by changing only the priors $\{p_i\}$. The Navier-Stokes application (github.com/RahulModak74/Navier-Stokes) demonstrates that the MW Framework naturally extends to Banach spaces for problems involving singularities and weak solutions, where classical Riemannian geometry is insufficient.

10. Related Work and Literature Review

10.1 Variational Autoencoders and Riemannian Geometry

The connection between VAE latent spaces and Riemannian geometry was pioneered by Arvanitidis et al. (2018), who showed that the decoder Jacobian induces a pullback metric and that geodesics in latent space correspond to meaningful interpolations. Hauberg (2019) further argued that 'only Bayes should learn a manifold', grounding the approach in Bayesian inference. The MW Framework goes beyond these works by: (i) incorporating **physics priors** directly into the latent space geometry (via W and the KL term); (ii) providing a **falsifiable regime** (Student-t with $\nu < 1$) that recovers quantum gravity phenomenology; (iii) proving an **$O(n)$ inference theorem** that is independent of the decoder's complexity.

10.2 Physics-Informed Machine Learning

Physics-informed neural networks (PINNs; Raissi et al., 2019) embed PDE residuals into the loss function. The MW Framework differs fundamentally: PINNs still require the PDE to be known in closed form and re-solve it for each new boundary condition. MW learns the solution *manifold* once; subsequent inference is a single forward pass. This distinction is analogous to the difference between 'simulating' and 'retrieving'.

10.3 Geometric Deep Learning

Bronstein et al. (2017, 2021) unified many geometric architectures (CNNs, GNNs, transformers) under a symmetry-based framework. The MW Framework is geometric in a different sense: the geometry is *learned* from unlabelled data under physical constraints, not imposed by the symmetry of a fixed domain.

10.4 Bayesian Nonparametrics and Manifold Learning

Gaussian process latent variable models (GP-LVM; Lawrence, 2005) and Bayesian manifold learning (Tosi et al., 2014) share the goal of representing data on a low-dimensional manifold. The MW Framework replaces the GP with a VAE for scalability (to $n=25,000$) and adds explicit physics priors to shape the manifold.

10.5 Quantum Gravity and Statistical Priors

The use of heavy-tailed priors to model Planck-scale fluctuations has precursors in the literature (Rovelli & Vidotto, 2015; Gambini & Pullin, 2011). However, no prior work has demonstrated a single Student-t($\nu < 1$) prior that yields five simultaneous emergent properties (Einstein residuals, unsupervised clustering of black-hole/wormhole/foam geometries, Hawking scaling, LIGO waveform match, Planck-scale discreteness). The MW Framework thus provides a new statistical paradigm for quantum gravity phenomenology.

10.6 Open Science and Reproducibility

All code is open-source and all synthetic data generators are included. The framework runs on a standard laptop, contrary to the HPC-dependence of traditional quantum gravity and fluid dynamics simulations. Three public repositories document the framework:

github.com/RahulModak74/mw-framework

github.com/RahulModak74/Navier-Stokes

github.com/RahulModak74/QUANTUM_GRAVITY_WITH_MW

A companion differential geometry primer (*Differential Geometry Required for the MW Framework*, Modak & Walawalkar, April 2026) provides the mathematical background needed to extend the framework to new domains.

11. Version 2.0 Preview: Manifold Algebra

Version 1.0 covers the single-manifold regime sufficient for the $10^6 \times$ theoretical maximum speedup against classical PDE solvers ($k=3$). Version 2.0 will introduce **manifold algebra** — four composition operations that build complex manifolds from tractable sub-manifolds, enabling the 10^{12} - 10^{20+} tier:

Manifold Algebra (Version 2.0 Preview). Given sub-manifolds $M_1, M_2, \dots, M_m$, four operations compose them:

Product: $M_1 \times M_2$ — independent systems, inference factorises, cost $O(m \times n)$.

Fibre Bundle: $M_1 \times_B M_2$ — M_2 geometry conditioned on M_1 state, hierarchical inference.

Gluing Map: $M_1 \cup_{\text{phi}} M_2$ — interface physics at shared boundary, cost $O(b)$, b = boundary dimension.

Marginalisation: integral_{M_j} — integrate out irrelevant sub-manifold, cost $O(1)$ amortised.

These operations will be detailed in a separate document.

12. Summary: The MW Framework in Five Steps

Step 1 — The Universe of Discourse

Choose n . Define R^n . Every conceivable state of the system lives here.

Step 2 — The Physically Admissible Region

Choose bounds a_i, b_i . Define $H = \text{product } [a_i, b_i] \text{ subset } R^n$. Physically impossible states are excluded.

Step 3 — Physics as Prior, Prior as Geometry

Write each governing physics law as a probability distribution p_i over H . Their intersection defines $M \text{ subset } H$. A physics-constrained VAE learns M from data + priors via ELBO maximisation.

Step 4 — Geometry Emerges from the Decoder

The learned decoder $D_{\text{psi}} : R^d \rightarrow H$ defines $M\text{-hat} = \text{Im}(D_{\text{psi}})$. The pullback metric $g = J_D^T W J_D$ equips $M\text{-hat}$ with Riemannian (or Lorentzian) geometry. All geometric objects — Christoffel symbols, curvature, geodesics, Van Vleck determinant — are computed automatically via autodiff. No hand calculation.

Step 5 — Universal $O(n)$ Inference

For any new observation x in H , compute $d_{\text{MW}}(x, M\text{-hat})$ in one forward pass — $O(n)$. The MW distance is a geometric surrogate for Synge's world function. Large d_{MW} means anomaly. Δ_{MW} gives uncertainty. Both scale to any dimension, any signature (Riemannian, Lorentzian, Banach), any physics domain.

Physics constrains solutions to a low-dimensional manifold; Bayesian inference on that manifold is always $O(n)$ — regardless of how hard the original equation was. Differential geometry is the language that makes this precise.

Acknowledgements

The authors thank the open-source Pyro and PyTorch communities for providing the probabilistic programming and automatic differentiation infrastructure that made this work possible. We also thank the anonymous curator of the 'Geometric Foundations of Quantum Theory and Gravity' Zenodo community for valuable feedback.

References

- [1] Milnor, J. (1965). *Topology from the Differentiable Viewpoint*. University Press of Virginia.
- [2] do Carmo, M.P. (1992). *Riemannian Geometry*. Birkhauser.
- [3] O'Neill, B. (1983). *Semi-Riemannian Geometry with Applications to Relativity*. Academic Press.
- [4] Synge, J.L. (1960). *Relativity: The General Theory*. North-Holland.
- [5] Kingma, D.P. & Welling, M. (2014). Auto-encoding variational Bayes. *Proc. ICLR*.
- [6] Bingham, E. et al. (2019). Pyro: Deep universal probabilistic programming. *JMLR*, 20(28).
- [7] Arvanitidis, G., Hansen, L.K., & Hauberg, S. (2018). Latent space oddity: on the curvature of deep generative models. *ICLR*.
- [8] Hauberg, S. (2019). Only Bayes should learn a manifold. *NeurIPS Workshop on Bayesian Deep Learning*.
- [9] Raissi, M., Perdikaris, P., & Karniadakis, G.E. (2019). Physics-informed neural networks. *J. Comput. Phys.*, 378, 686-707.
- [10] Bronstein, M.M., Bruna, J., LeCun, Y., Szlam, A., & Vandergheynst, P. (2017). Geometric deep learning. *IEEE Signal Process. Mag.*, 34(4), 18-42.
- [11] Lawrence, N.D. (2005). Probabilistic non-linear PCA with Gaussian process latent variable models. *JMLR*, 6, 1783-1816.
- [12] Tosi, A., Hauberg, S., Vellido, A., & Lawrence, N.D. (2014). Metrics for probabilistic geometries. *UAI*.
- [13] Dai, B. & Wipf, D. (2019). Diagnosing and enhancing VAE models. *ICLR*.
- [14] Gal, Y. & Ghahramani, Z. (2016). Dropout as a Bayesian approximation. *ICML*.
- [15] Rovelli, C. & Vidotto, F. (2015). *Covariant Loop Quantum Gravity*. Cambridge University Press.
- [16] Gambini, R. & Pullin, J. (2011). *A First Course in Loop Quantum Gravity*. Oxford University Press.
- [17] Nash, J. (1956). The imbedding problem for Riemannian manifolds. *Annals of Mathematics*, 63(1), 20-63.
- [18] Modak, R. & Walawalkar, R. (2026). MW Framework: Complexity-Theoretic Alternative to PDE Solvers. Zenodo. DOI: 10.5281/zenodo.15304813.
- [19] Modak, R. & Walawalkar, R. (2026). A Student-t($\nu=0.8$) Prior on Measurement Processes Yields a Riemannian Manifold Simultaneously Consistent with GR and QM. Zenodo. DOI: 10.5281/zenodo.19304813.
- [20] Modak, R. & Walawalkar, R. (2026). Differential Geometry Required for the MW Framework. Bayesian Cybersecurity Pvt Ltd, April 2026.

Rahul Modak — Co-developer, MW Framework | Bayesian Cybersecurity Pvt Ltd, Thane / Mumbai

Dr. Rahul Walawalkar — Co-developer, MW Framework | Carnegie Mellon University; NETRA; Caret Capital